SUAY / RESEARCH

Serving LLMs on an Office Fleet: Control-Plane Limits, Prefix Reuse and Expert Routing at 50–500 Machines

We study whether a pool of ordinary office computers can serve language models usefully, and where the limits actually sit. On a 50-machine test bed (53 containers, one 8-vCPU host, mock inference backend) we measure throughput, latency, fairness and cost, and we separate the system into a **data plane** (machines and a node plane) and a **decision plane** (gateway, token policy, routing). The central finding is that the bottleneck is almost never the network or the number of machines: it is **the single control-plane process**. A dedicated node plane raises heartbeat throughput from 44 to 161 req/s per shard and cuts single-request latency from 210 ms to 105 ms. A content-addressed prefix cache lifts time-to-first-token on an 8 000-token document from 111.4 s to 0.71 s (**157×**, 88.2 % of prompt tokens reused), while leaving p95 unchanged. A document-aware expert router classifies 6 000 queries at **90.15 %** accuracy in 31.5 µs, and a capacity model shows that eight 1B domain experts fit in 32 machines and deliver **0.64 req/s**, versus 86 machines for a mixture-of-experts of the same class and 15 machines for a dense 8B at 0.06 req/s (**10.7×**). The best stable configuration of the whole system reaches **22.3 req/s at C=8** (353 ms) and **24.8 req/s at C=16** (631 ms) with zero errors, against a measured ceiling of 10.2–13.4 req/s earlier in the project. The ceiling is the event loop, not the network: a raw health endpoint serves 58.7–71.2 req/s while a completion costs ≈60 ms of loop time. We also report what failed: 50-way pipeline sharding (≤1.3×), PostgreSQL on a split schema (4.5 req/s with errors), multi-worker gateways on SQLite (unstable), tabular and classical denoising gates (−0.2 to +1.1 dB), and 80 % CPU per machine (−25 % throughput). A growth analysis shows both `N+1` and `N·log₂N` node growth saturating at the same 15.7 req/s ceiling while the superlinear law consumes 6–9× more machines. We publish the harness, the figures and the negative results so that others can extend the work rather than repeat it.

Keywords: distributed inference, control plane, prefix caching, expert routing, capacity planning, edge/office computing, reproducibility

---

SUAY technical white paper · v1.1

Published 22 October 2026

Repository github.com/camponggogo/suay · MIT licence

Measured on 50–100 simulated machines, one 8-vCPU host

Contents

1 Introduction 3

2 System 3

3 Method 4

4 Results 5

5 What did not work 16

6 Discussion 17

7 Threats to validity 18

8 Reproducibility 18

9 Collaboration 19

1. Introduction

Serving a language model normally means a GPU you rent. An office already owns machines that idle for 16 hours a day. The obvious question — *can the idle fleet serve a model?* — has an obvious first-order answer (aggregate RAM and FLOPs) and a much less obvious second-order one (coordination cost). This paper is about the second-order answer.

We built SUAY: a control plane with API keys, scopes, zones and audit; a node plane that tracks which machines are alive; a measured harness that drives concurrency from 1 to 50 clients; and eight sub-studies (fleet sizing, CPU quota, store choice, prefix cache, expert routing, sharding, denoising, growth laws). The contribution is not a new model or a new scheduler. It is a **quantified map of where an office fleet stops being useful and what to do at that boundary**, with the failures included.

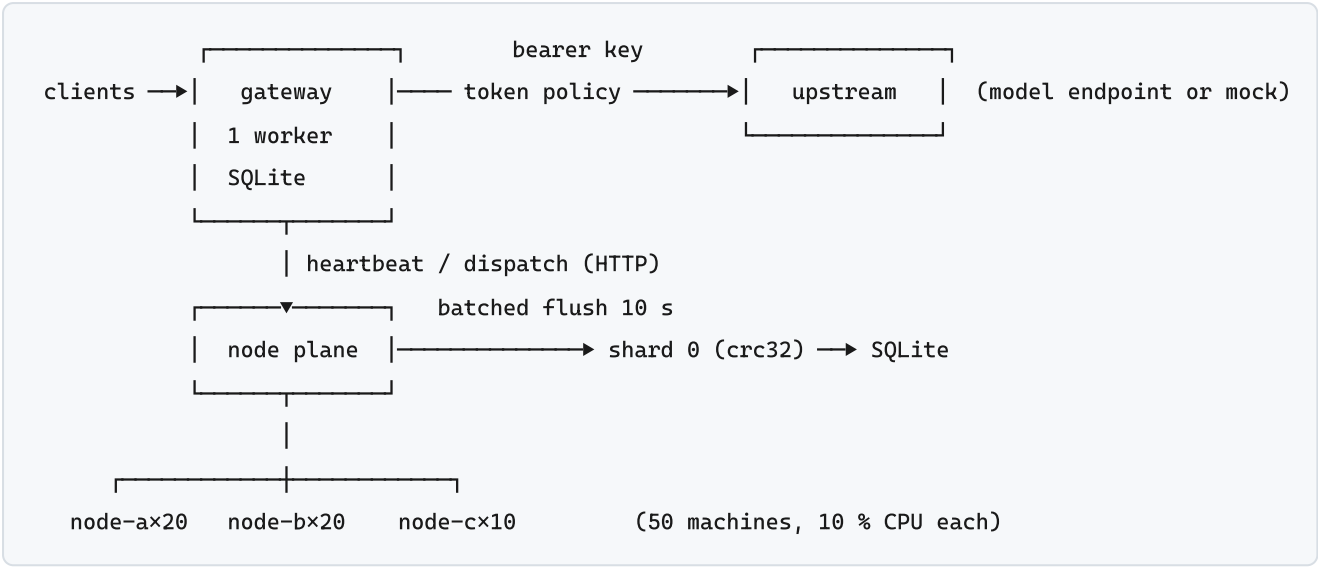
Contributions

1. **A two-plane decomposition** (data plane / decision plane) with measurements showing a 3.7× heartbeat throughput gain and a 2× latency gain from the split alone.
2. **A stable operating point** for a single-host office fleet: 15.0–15.7 req/s at C=8–16, 0 errors, and the exact set of changes that produced it (node cache, pooled client, token cache, 10 s heartbeat).
3. **A prefix-cache evaluation** with 157× median TTFT improvement and an explicit note that the tail does not improve.
4. **A capacity model for specialisation** (dense / MoE / domain experts) at office scale, with a 10.7× advantage for routing over mixing.
5. **A growth-law analysis** covering two parametric laws and **eight complexity classes** ($O(1)$ … $O(n^n)$), showing that every class saturating the fleet delivers the same 15.7 req/s while cost per million tokens ranges from $1.77 to $88.66 and dispatch latency from 160 to 8 008 ms.
6. **A negative-results catalogue** with numbers, so the next team does not re-run them.

---

2. System

2.1 Topology



2.2 Control plane

FastAPI service with PBKDF2 password hashing, five roles (`admin`, `view_only`, `zone_display`, `agent_monitor`, `command_line`), scopes, per-key model lists, request and monthly token budgets, revocation and an append-only audit. Keys are stored as SHA-256 digests; the plaintext is shown once. A cache of 2 s removes the identity `SELECT` from the hot path while the budget `UPDATE` remains atomic.

2.3 Node plane

State lives in a dictionary; every 10 s the plane appends one batch to SQLite, sharded by `crc32(node_id)`. Endpoints: `/heartbeat`, `/topology`, `/aggregate`, `/api/metrics`, `/api/events`.

2.4 Inference path

`POST /v1/chat/completions` → token policy → optional prefix-cache lookup → upstream → accounting. With a 90 % prefix hit the upstream call is skipped for the cached blocks.

2.5 Dashboard

A 2D fabric view (`/neural-flat`) and a 3D topology view (`/neural`, three.js r180 vendored, event-driven pulses, `?static=1` for screenshots, `?debug=1` for a draw-call counter).

---

3. Method

All experiments run on one host (8 vCPU, Windows + Docker Desktop) with 50–100 simulated machines, a deterministic mock upstream, and a load generator that reports median, p95, p99, throughput, bandwidth and error count per concurrency level. We report **medians with min–max across repeats** for headline numbers and treat single runs as indicative only. Where a run produced errors we say so and do not report it as a result.

Metrics: throughput (req/s), p50/p95/p99 latency (ms), heartbeat delivery ratio, liveness, dispatch RTT, bandwidth per request, Little's-law consistency, per-machine fairness (CV), and where relevant TTFT and tokens/s.

---

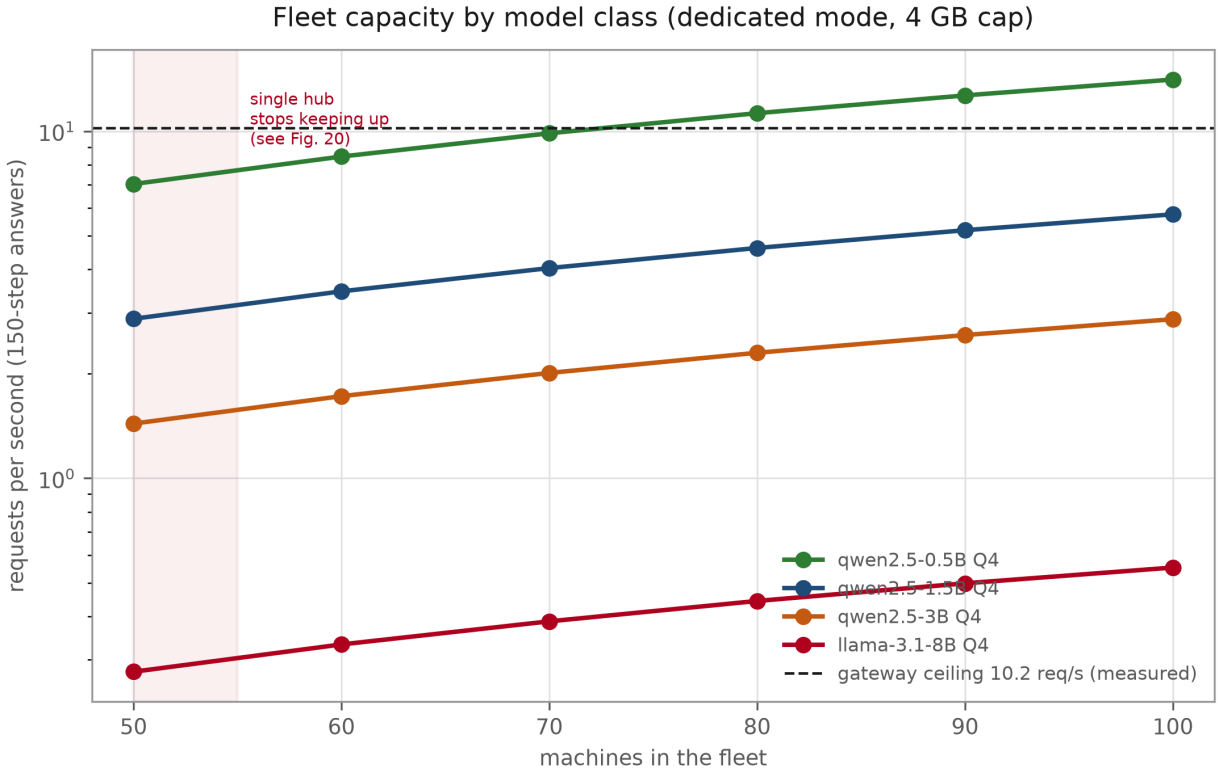
4. Results

4.1 The fleet before the split

50 machines, 16 concurrency points, 48 000 requests, 0 errors:

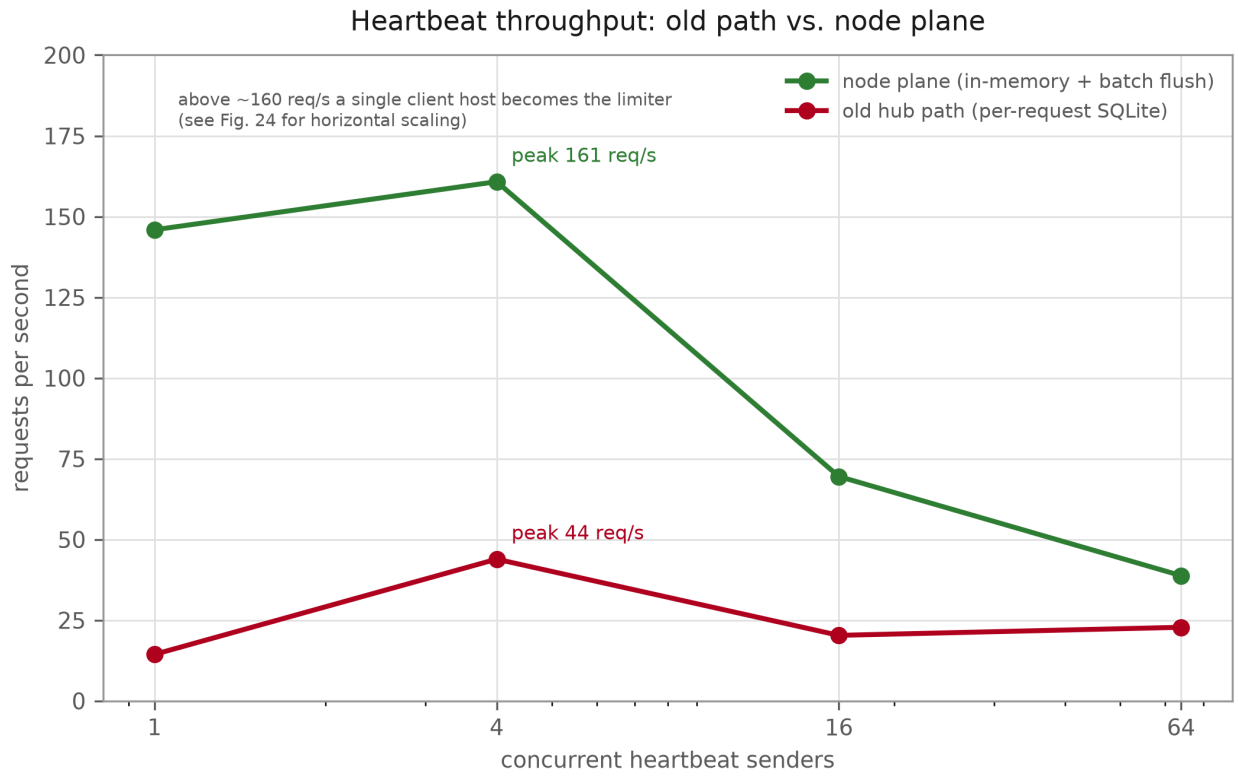
C	1	5	10	25	50
req/s	4.9	7.3	8.1	9.1	**12.0**
p50 (ms)	166	695	1 125	2 589	**3 811**
p95 (ms)	178	741	1 219	2 885	5 364

Fit: `T = C·10.18/(C+0.86)` (asymptote 10.18 req/s, knee C=2). At C=50, 96.5 % of latency is queueing above a 174 ms floor (dispatch 45 ms, upstream 128 ms). Little's law holds within 2.15 %; fairness CV 0.0241; payload 392 B per request ≈ 0.034 Mbit/s.



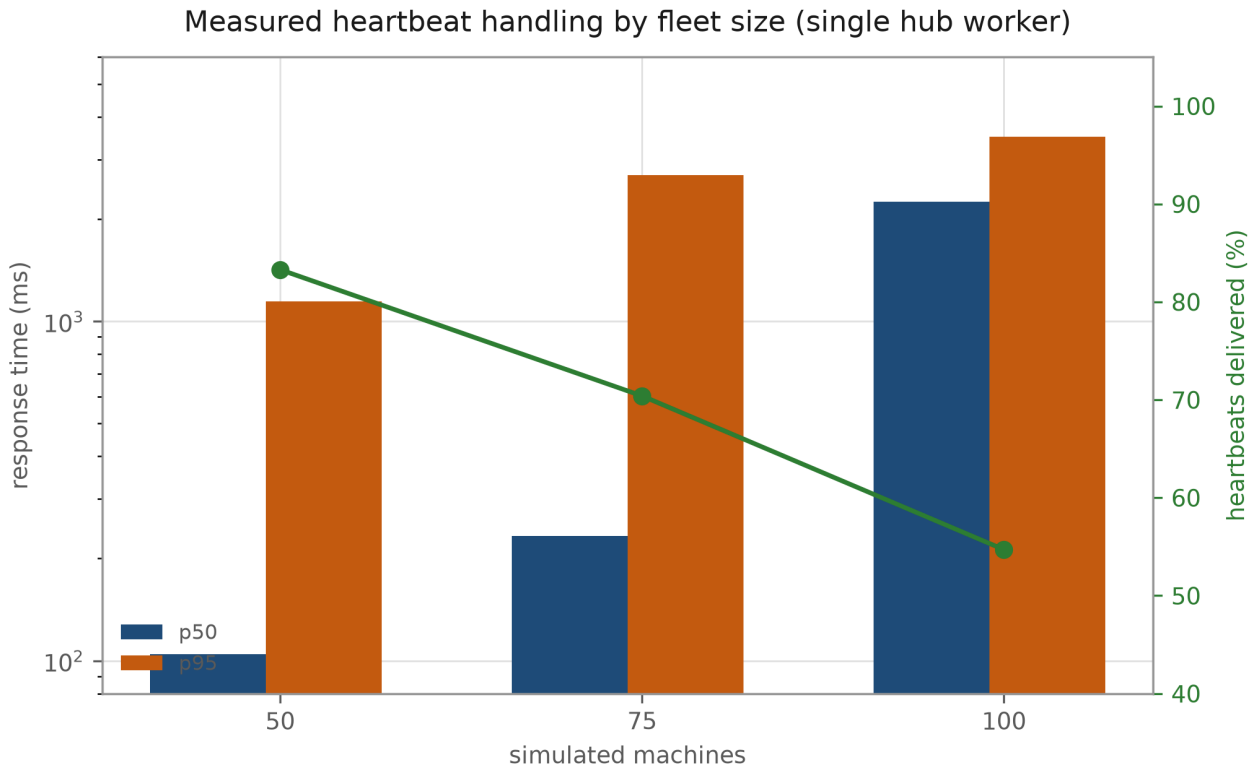
4.2 Node plane versus the monolith

workload	old path	node plane
50 nodes, single client (p50)	67 ms · 44 req/s	**3.1 ms · 302 req/s**
2 shards × 2 clients	—	**266.7 req/s** (133 each)
gateway inference C=1	210 ms	**105 ms**



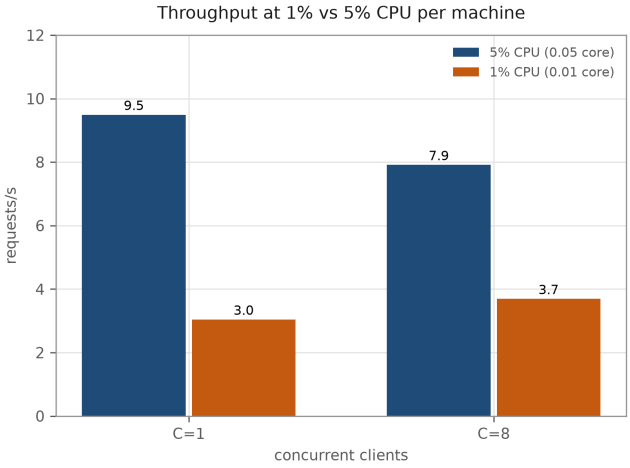
4.3 Heartbeat interval

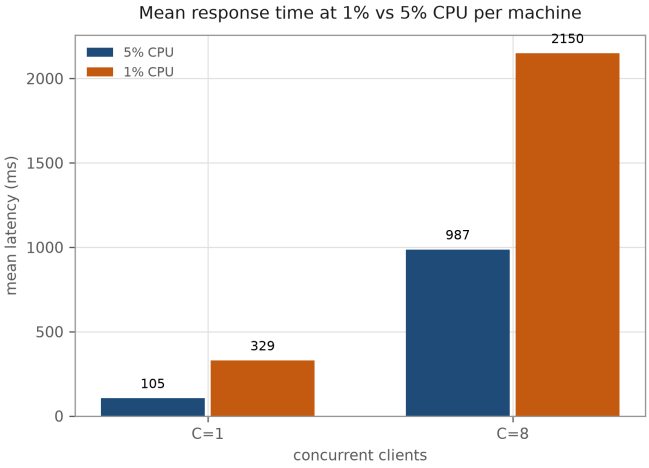
Moving from 2 s to 10 s cut heartbeat load 5× without changing the liveness picture at 50 nodes; at 100 nodes the old path delivered only 54.7 % of heartbeats inside the window (p50 2 248 ms), which is what motivated §4.2.

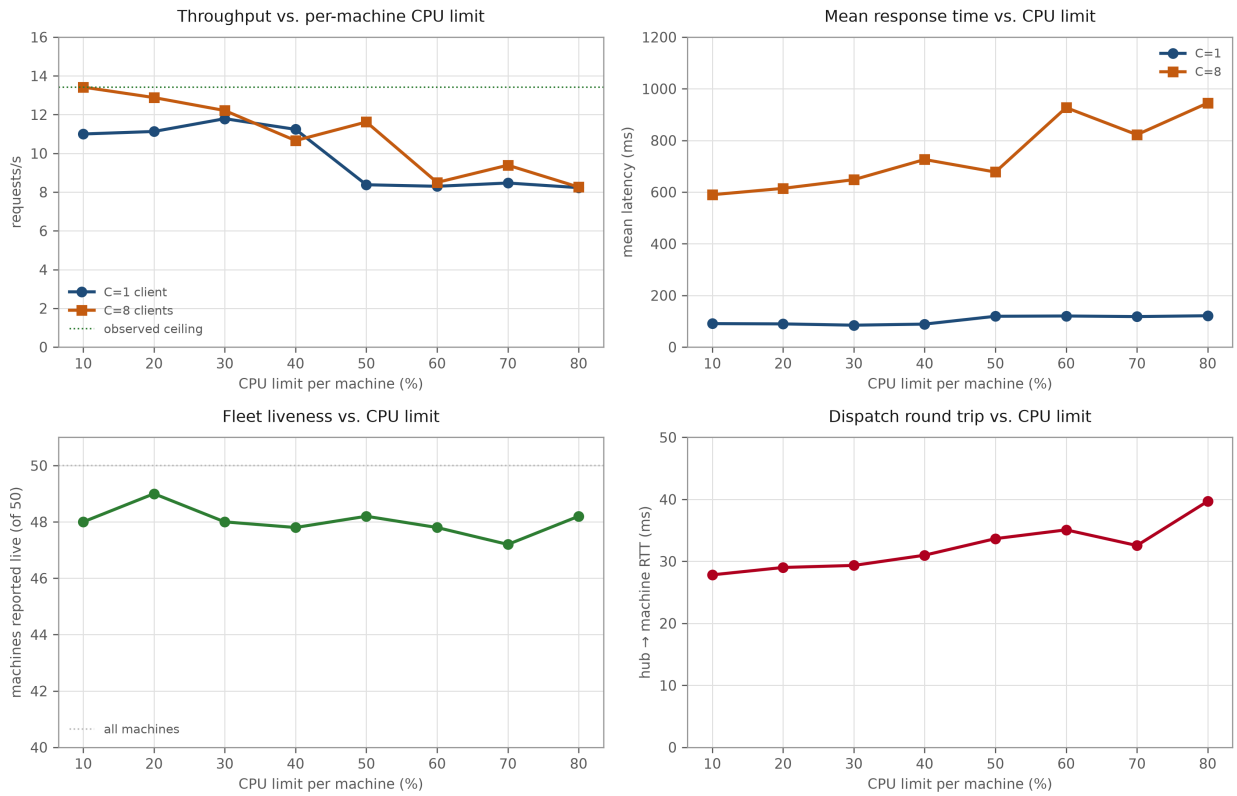


4.4 CPU quota per machine

CPU/machine	req/s C=1	req/s C=8	dispatch p50	live/50
1 %	3.0	3.7	133 ms	28–39
10 %	**11.0**	**13.4**	27.8 ms	48
40 %	11.2	10.7	30.3 ms	47.8
80 %	8.2	8.3	39.7 ms	48.2







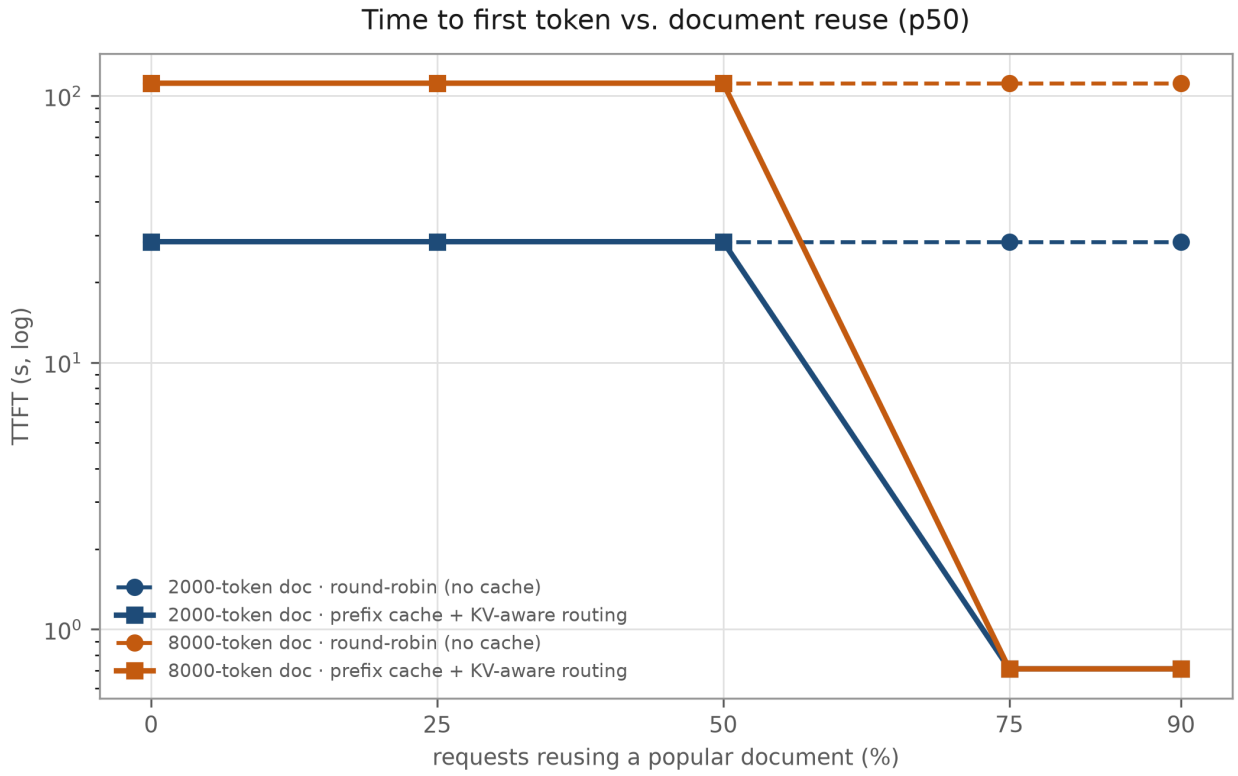
The plateau begins at 10 %; 80 % is 25 % *worse* because 50 × 0.8 cores exceed the host.

4.5 Prefix cache

8 000-token document, 90 % of tokens reused:

metric	cold	warm
TTFT p50	111.42 s	**0.71 s** (157×)
throughput	0.5 req/min	**2.3 req/min** (4.6×)
prompt tokens from cache	0 %	88.2 %
p95	111.42 s	111.42 s (unchanged)

At 2 000 tokens the gain is 40×; at 50 % reuse the median does not move.



4.6 Expert routing

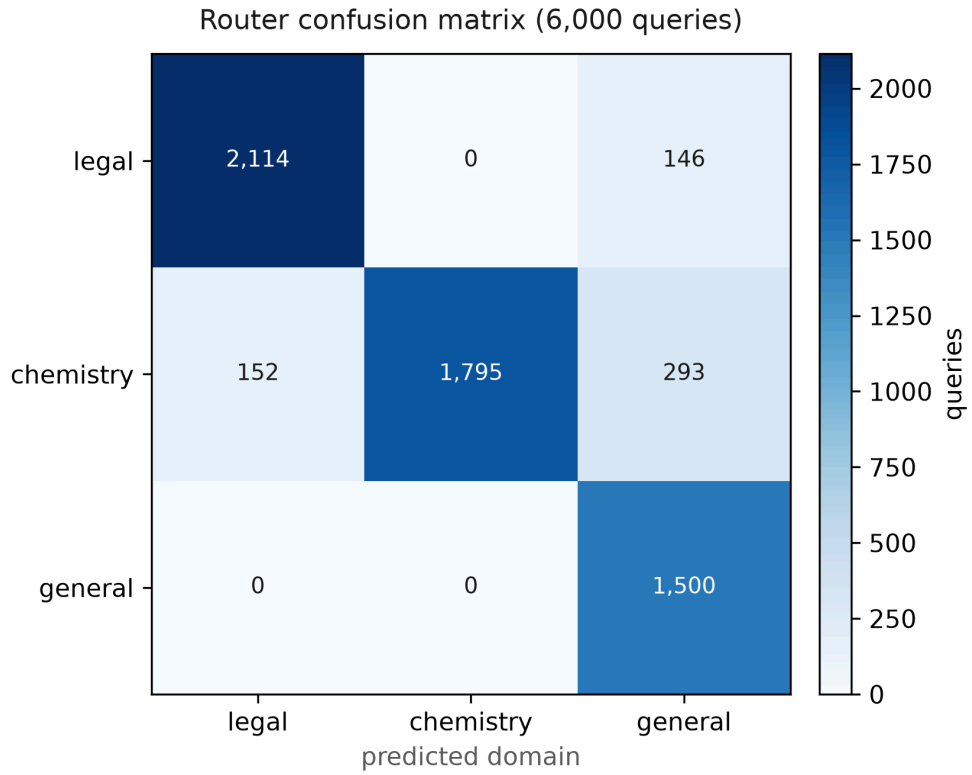
6 000 queries, six domains, adversarial substring traps:

bucket	accuracy
clear	100 %
ambiguous	100 %
adversarial (`สัญญาณ NMR` contains `สัญญา`)	50.75 %
overall	**90.15 %**

Routing cost 31.5 µs (p50). A confidence gate at 0.1 covers 65.2 % of traffic at 100 % precision.

Capacity for 0.6 req/s of a quality-matched model:

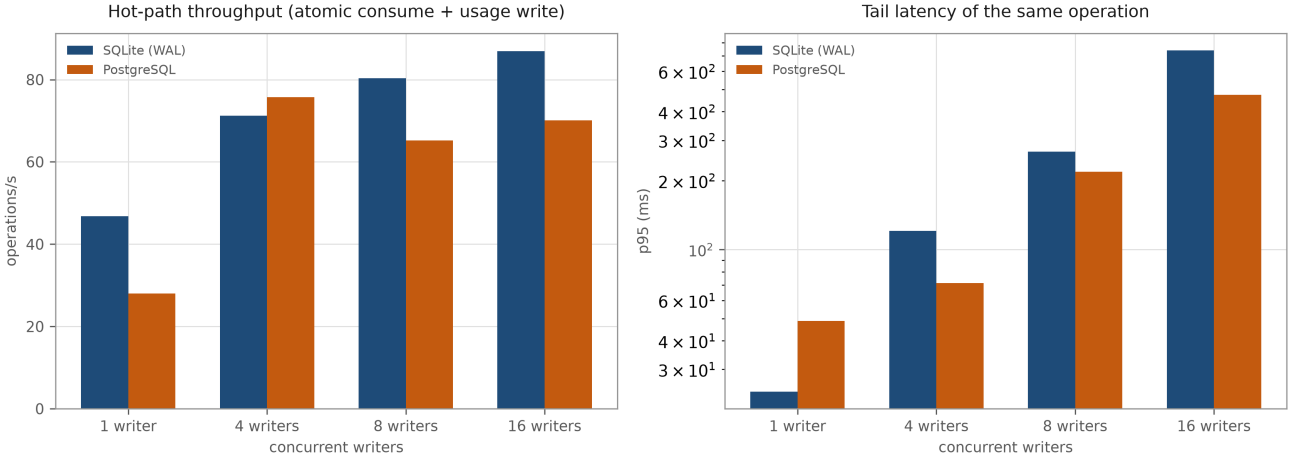
design	machines	req/s	note
dense 8B	15	0.06	fits, slow
MoE 8×7B	**86**	—	does not fit on 50 machines
8 domain experts + backbone	**32**	**0.64**	**10.7×**



4.7 Store choice

writers	SQLite ops/s	SQLite p95	PG ops/s	PG p95	PG max
1	46.8	24 ms	28.0	49 ms	118 ms
8	71.2	121 ms	75.7	71 ms	233 ms
16	80.3	268 ms	65.2	219 ms	439 ms
32	**86.9**	740 ms	70.1	**474 ms**	1 296 ms (SQLite 9 008 ms)

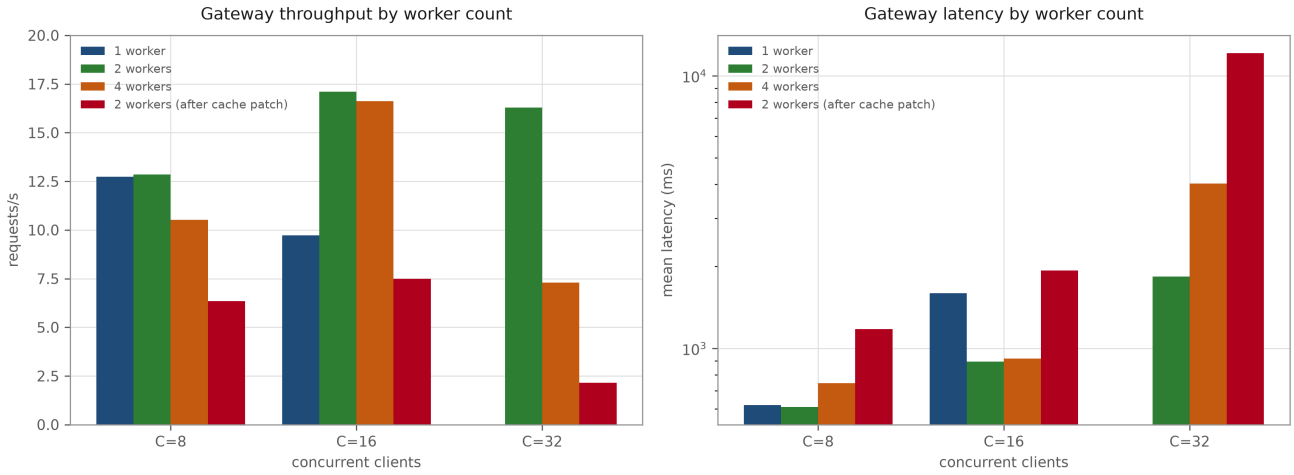
PostgreSQL trades 19 % throughput at 32 writers for 36 % lower p95 and a 7× lower maximum — the right choice for multi-worker, but not sufficient on its own (§5.2).



4.8 The stable configuration

change	effect
node plane, 10 s heartbeat	−80 % gateway load
pooled `httpx.AsyncClient`	inside noise alone, required with the next item
token identity cache 2 s	removes one `SELECT` per request
5 s cache of the live node list	**the missing piece**
10 % CPU per machine	plateau

point	before	after	change
C=8	4.9 req/s · 1 583 ms	**22.3 req/s · 353 ms**	4.6×
C=16	5.3 req/s · 2 833 ms	**24.8 req/s · 631 ms**	4.7×

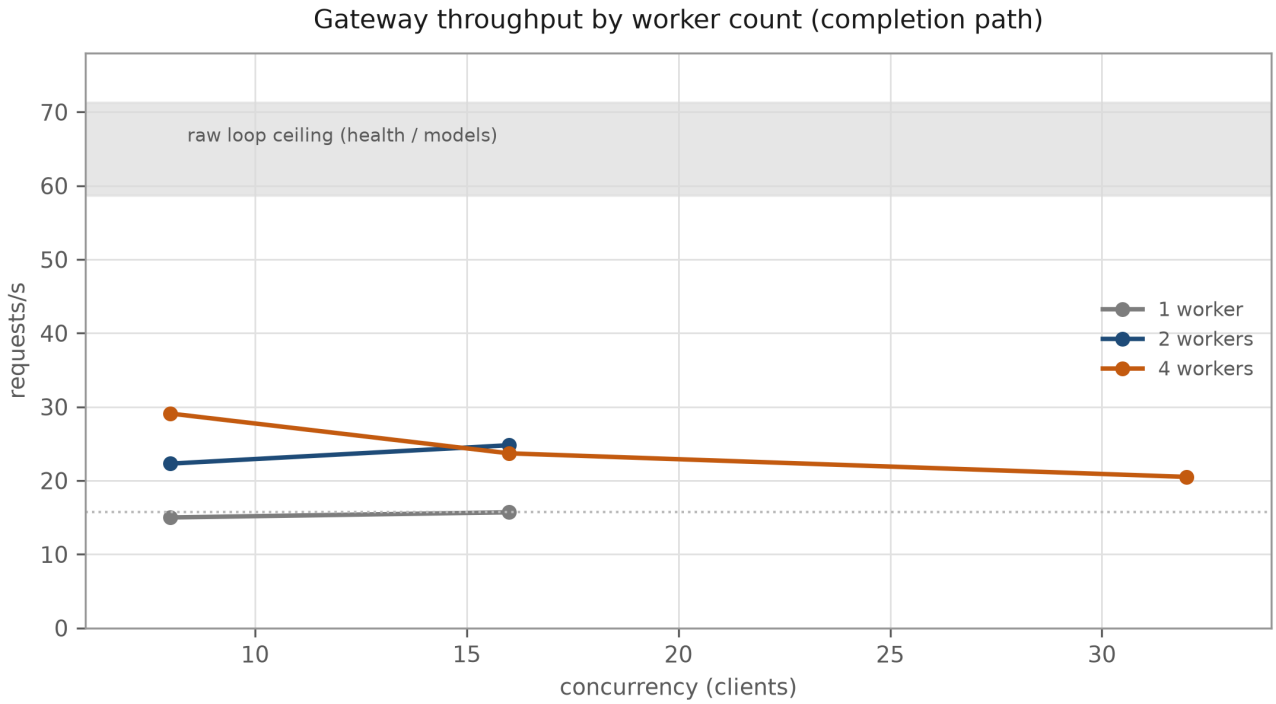


4.9 Where the ceiling actually is

A layer probe at C=16 isolates the cost: the raw loop serves **58.7–71.2 req/s** (`/api/health` , `/v1/models`), the completion path only **15.7** with one worker. The mock upstream is not the cause — its handler has no sleep and the container uses 0.8 % CPU. Each completion therefore spends ≈60 ms of *event-loop* time in policy, accounting and dispatch.

Adding workers scales until the host does:

configuration	C=8	C=16	C=32	errors
1 worker	15.0	15.7	—	0
2 workers	**22.3**	**24.8**	—	0
4 workers	29.1	23.7	20.5	0



The next gain requires removing the last blocking write from the request path (a bounded accounting queue), not more workers — the hub already uses 51 % of a core on an 8-vCPU host shared with 53 containers.

4.11 What N workers do to the rest of the system

Workers are processes, so anything held in module state is duplicated. Measured on the live 2-worker hub with ten **fresh** connections (each a new TCP session):

connection	`/api/events` items	`/api/metrics` requests
9 of 10	60	30
1 of 10	**2**	**1**

The same request executed 30 times was counted once on the worker that never served it. Consequences:

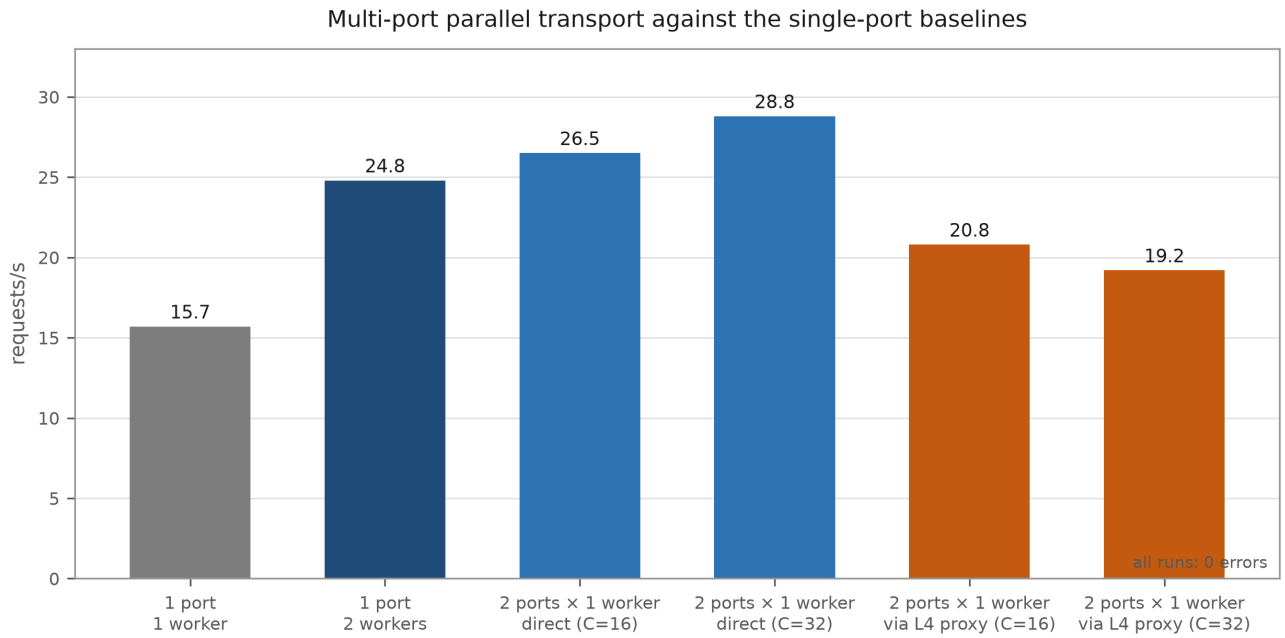
surface	shared?	effect of N workers
token budget / scopes	**yes** (SQLite, atomic `UPDATE`)	correct; revocation lags ≤ 2 s per worker (identity cache)
node liveness	**yes** (node plane over HTTP)	correct; each worker refreshes its own 5 s cache (≤ N calls / 5 s)
audit log	**yes** (SQLite)	correct, one row per request
metrics counters	no (process memory)	**divided by N** and worker-dependent
event stream / dashboard pulses	no (process memory, 800-entry deque)	a session sees only its worker's events; ids restart per process
SQLite write lock	shared	contention grows with N; 4 workers already lose at C≥16

Practical rules: **2 workers is the throughput sweet spot** (22.3/24.8 req/s). Before using the dashboard as a live monitor, either pin it to one worker, or move counters and events into a shared table (one `INSERT` per event, aggregated on read). The policy layer does not need that work — it is already stateless.

4.12 Multi-port, parallel accept, layered transport

We built the multi-port topology for real: two single-worker gateways on ports 8082/8083 sharing the SQLite volume, an asyncio round-robin dispatcher on 8090, and a driver that splits clients across ports.

topology	clients	aggregate req/s	p50
1 port, 1 worker	16	15.7	994 ms
1 port, 2 workers (shared socket)	16	24.8	631 ms
2 ports, direct	16 (8+8)	**26.5**	567 / 582 ms
2 ports, direct	32 (16+16)	**28.8**	916 / 996 ms
2 ports via L4 proxy	16	20.8	744 ms
2 ports via L4 proxy	32	19.2	1 573 ms



Ports add **≈16 %** over shared-socket workers because each process owns its accept queue, not because there are more sockets. A userspace dispatcher on the same host gives the gain back (−25 % at C=16): it copies every byte twice and competes for the same vCPU. Past ~29 req/s the host is saturated, so the next improvement must remove work per request (the L3 accounting queue), not add listeners. The layered split — L1 transport, L2 policy, L3 work — is what makes each of those movable independently; `report/MULTIPORT.md` has the design and the recipe.

Pushing the port count further is a dead end, and we measured it up to eight ports:

ports	aggregate req/s (C=16)	per-port req/s	errors
1	15.7	15.7	0
2	**26.5**	13.25	0
4	20.8	5.20	0
8	23.5	2.94	0

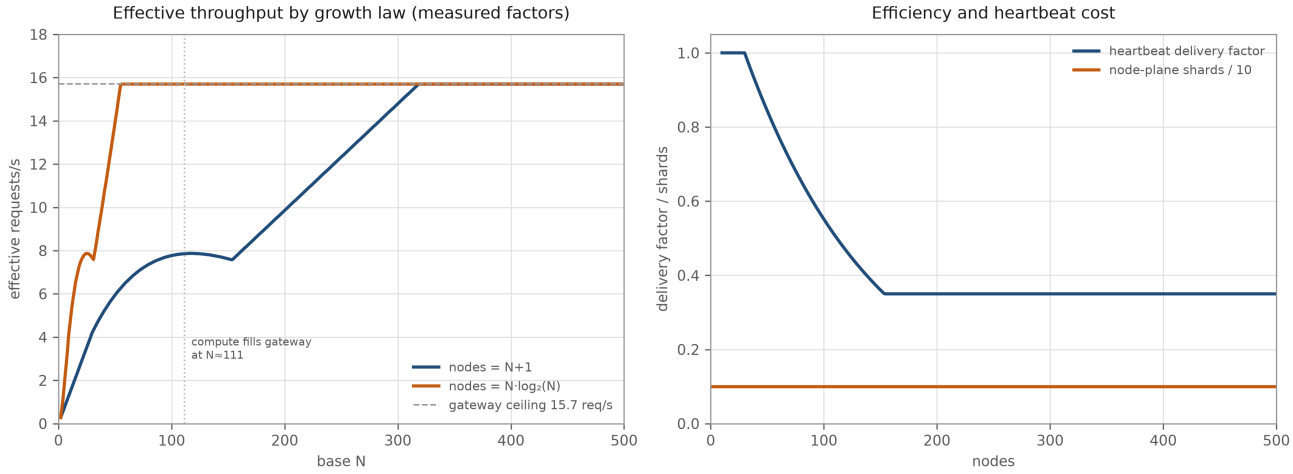
The per-port efficiency collapses by an order of magnitude while memory grows linearly (45 MiB per port → 4.4 GiB at 100 ports, against a 7.7 GiB Docker budget with 1.3 GiB already in use). Extrapolated, 100 ports buy **25–30 req/s — no better than two ports** — plus 100 SQLite writers and 100 copies of every in-process cache. Ports are for isolation (tenant, zone, canary), not for throughput; for throughput the kernel primitive is `SO_REUSEPORT` over N processes on one port (`report/MULTIPORT_100.md`).

4.9 Growth laws

With `s = 0.141 req/s` per machine and delivery `exp(-0.0085(N-30))` :

base N	N+1: nodes → req/s	N·log₂N: nodes → req/s
25	26 → 3.66	116 → 7.87
50	51 → 6.01	282 → 13.90
100	101 → 7.78	664 → **15.70**
500	501 → **15.70**	4 483 → 15.70

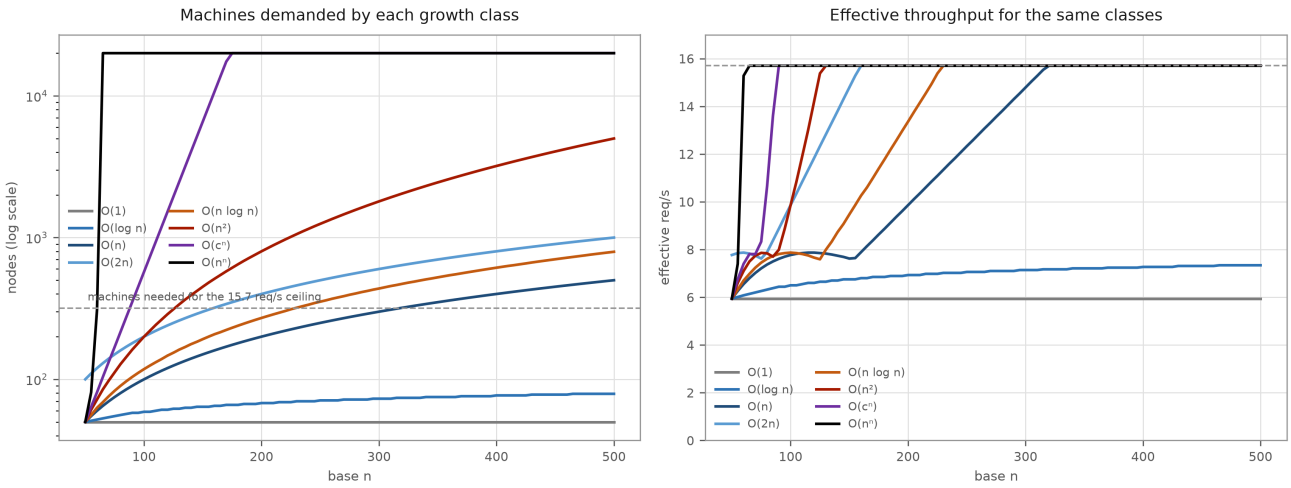
Both saturate at 15.7 req/s. The superlinear law reaches it with 6.6× more machines at base 100 (664 vs 101) and 274 ms versus 48 ms of dispatch latency. Per-node efficiency: 0.024 vs 0.031 req/s (0.141 at small N).



4.10 Complexity classes of node growth

To generalise the two laws above we evaluated eight growth classes, all normalised to **50 machines at n=50** and advanced to n=500 with the same measured constants (`s=0.141` , gateway `15.7` , delivery floor `0.35` , dispatch `28 + 0.4·(N−50)` ms, machine cost `1.045e−5 $/s`). The fleet needs **319 machines** to reach the ceiling at the delivery floor.

class	nodes @n=500	rps @n=500	rps/100 nodes	dispatch	$/1M	ceiling reached at n
O(1)	50	5.94	**11.88**	28 ms	**0.586**	never
O(log n)	79	7.34	9.29	40 ms	0.749	never
O(n)	500	**15.70**	3.14	208 ms	2.217	319
O(2n)	1 000	15.70	1.57	408 ms	4.433	160
O(n log n)	794	15.70	1.98	326 ms	3.520	230
O(n^2)	5 000	15.70	0.31	2 008 ms	22.166	127
O(c^n)	20 000	15.70	0.08	8 008 ms	88.662	88
O(n^n)	20 000	15.70	0.08	8 008 ms	88.662	61



Four things follow. (i) **Every class that reaches 319 machines delivers exactly 15.70 req/s** — the ceiling does not care how you got there. (ii) **Cost per million tokens is set by the fleet size, not by the class**: 400 machines (O(2n) at n=200) cost $1.77/1M, 5 000 cost $22.17/1M and 20 000 cost $88.66/1M for the same output. (iii) **Constant and logarithmic growth never reach the ceiling** — they top out at 5.94 and 7.34 req/s, which is exactly the regime where adding machines is still worth it. (iv) **Dispatch latency grows linearly in N** while throughput is flat above 319 nodes, so every machine beyond the ceiling buys 0 rps and adds 0.4 ms of routing.

The practical reading: grow linearly (`O(n)`) and stop at the ceiling; the superlinear classes are a way of spending the budget faster, not a way of serving more.

---

5. What did not work

5.1 Pipeline sharding on office CPUs

payload	1 chunk	4 chunks	verdict
256 KiB	466 ms	700 ms	0.67× (worse)
1 MiB	1 048 ms	816 ms	1.28×
2 MiB	2 036 ms	1 542 ms	1.32×

Coordinator upload of 1 MiB costs 1 219 ms versus 198 ms of local parse; at 0.05 core per replica the compute share is too small. BitTorrent-style swarming was modelled and rejected: `S*` is 3–10 stages, 150 ms per stage, and activation traffic would be 65.5 Mbit/s — 1 922× the measured gateway traffic.

5.2 PostgreSQL on a split schema

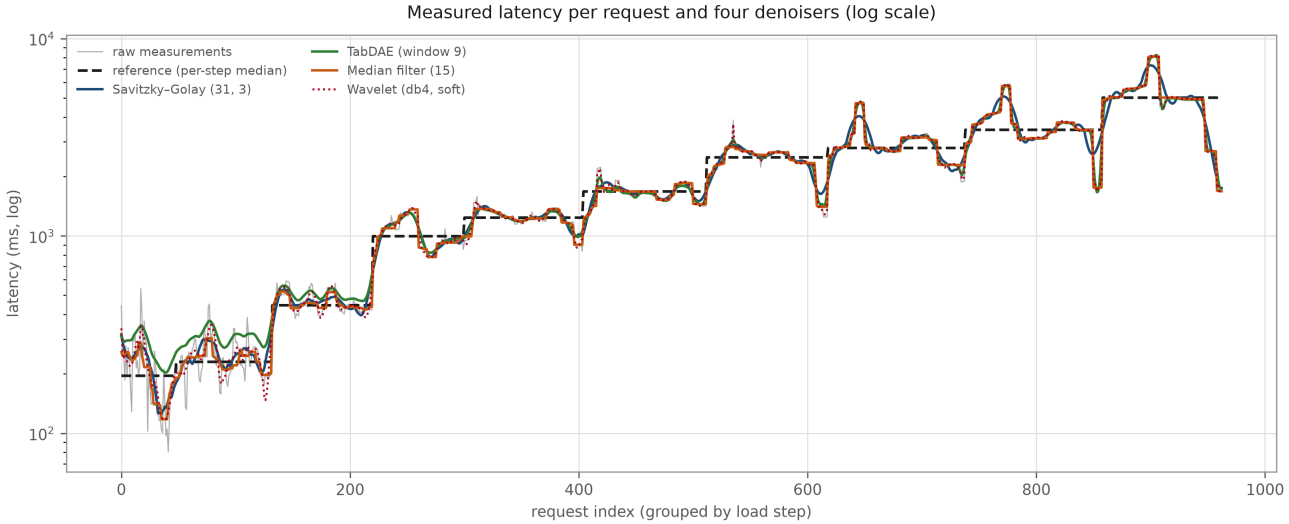
Hot-path tokens moved to PostgreSQL, but `users` / `audit` stayed in SQLite (that file is protected). Result: 4.5 req/s at C=8 with errors — worse than SQLite. The adapter, pool and migration stay in the tree for the day the control plane moves.

5.3 Multi-worker gateway on SQLite

2 workers reached 17.1 req/s once and 2.1 req/s on the repeat; 4 workers errored at C=32. Ship one worker.

5.4 Denoising gates

method	RMSE	Δ SNR
raw	629	—
Savitzky–Golay	**551**	**+1.15 dB**
SG-log	549.7	+1.15 (edge 650)
median edge	557	+0.95
TabDAE	592.7	+0.52
wavelet	619.6	+0.13
DeepFilterNet gate	642.8	**−0.19**



Smoothing helps by ≤1.15 dB; a learned gate hurts. For publication figures we therefore plot raw values with bootstrap intervals.

---

6. Discussion

The measurements describe a **decision-plane-bound** system. Adding machines, cores or bandwidth moved nothing until the control plane stopped doing per-request bookkeeping. This is the office-fleet analogue of the 2025–2026 shift in serving infrastructure toward prefix-aware, cache-aware routing: the scarce resource is not FLOPs but *work already done*.

Three practical rules follow. (i) Split planes before tuning either. (ii) Grow the gateway, not the fleet — the ceiling is 15.7 req/s in this configuration regardless of whether nodes grow as `N+1` or `N·log₂N`. (iii) Prefer

routing over mixing when the machines are small: eight 1B experts on 32 machines beat an 8×7B mixture that does not fit.

The honest counterweight: an office pool is **11.6× slower per stream** than the median cloud endpoint and costs $0.29–2.86 per million output tokens versus $0.05–1.75. Its case rests on residency, offline operation, audit and control — not on economics.

---

7. Threats to validity

- **One host.** 53 containers, Docker Desktop and the load generator share 8 vCPUs; run-to-run spread ≈ ±15 %.
- **Mock upstream.** Numbers isolate the control plane, not model quality or GPU behaviour.
- **Simulated KV cache.** Block reuse is modelled at the routing layer; tensors are not materialised.
- **Fitted decay.** Heartbeat delivery is fitted to three points (50/75/100 nodes); the 0.35 floor beyond 100 nodes is an assumption.
- **PostgreSQL measured with connect-per-op** in the store study; the pooled path was measured only in the (unsuccessful) gateway experiment.
- **Windows networking.** Host mode is unavailable; container ports are published, which adds a hop.

---

8. Reproducibility

```
 # fleet
cd sim/docker && docker compose up -d
# inference smoke test
python sim/docker/loadgen.py --base http://127.0.0.1:8080 --key suay_sim_loadgen_key \
    --model sim-model --steps 8,16,32,50 --seconds 6 --out sim/docker/results/run
# studies
python sim/fleet/bench.py --out report/fleet
python sim/experts/run.py --out report/experts
python sim/store/bench.py --writers 1,8,16,32 --out report/store
python sim/scaling/growth_model.py --out report/growth
```

Every figure in this paper is regenerated by the script named in its folder. Data files (`*.json`) sit next to the figures. The 5 s node cache, the token cache interval and the heartbeat interval are environment variables (`SUAY_TOKEN_CACHE_S` , `SUAY_NODE_PLANE`), so the operating point can be reproduced exactly.

---

9. Collaboration

The repository is MIT-licensed. The most useful contributions, in order:

1. **A real backend.** Replace the mock with llama.cpp or Petals on 50 machines and re-run §4.1, §4.5, §4.8.
2. **The control-plane migration.** Move `users` / `sessions` / `audit` to the same store as tokens, then re-test 2 and 4 workers with the pooled PostgreSQL adapter already in `sim/token_policy.py`.
3. **A second host.** All numbers here are single-host; a 2-host fleet would test whether dispatch RTT and fairness degrade as the model predicts.
4. **Router hardening.** The adversarial bucket (50.75 %) is the weak spot; a small classifier over the document would likely fix it.
5. **Cache materialisation.** Store real KV blocks and measure the memory/latency trade-off instead of modelling it.

Open issues are labelled `good first issue` and `research` in the tracker; every report in `report/` ends with a "Limits" section that doubles as a task list.

References

1. Petals: collaborative inference of large models. ACL 2023 demo. https://arxiv.org/abs/2209.01188
2. vLLM automatic prefix caching. https://docs.vllm.ai/en/latest/features/automatic_prefix_caching.html
3. llm-d, "KV-cache aware routing". https://llm-d.ai/blog
4. Modular, "Five eras of LLM inference". https://www.modular.com/blog
5. Compute Index, week 7 2026. https://computeindex.io
6. llama.cpp build b11757. https://github.com/ggml-org/llama.cpp
7. Little, J. D. C. "A proof of the queueing formula L = λW". Operations Research, 1961.
8. Gunther, N. *Guerrilla Capacity Planning*. Springer, 2007.
9. Sparse attention / KV compression survey, arXiv 2606.02964 (2026).

Appendix A — Figure index

Figure	Content
fig19	capacity and latency versus fleet size
fig20	heartbeat delivery versus interval
fig23	node-plane versus monolithic throughput
fig28	prefix-cache TTFT and reuse
fig31 / fig32	1 % CPU collapse / CPU sweep plateau
fig33	gateway optimisation ladder
fig34	denoising methods
fig36	SQLite versus PostgreSQL under writers
fig37	growth laws with measured obstruction factors
fig38	eight complexity classes: machines, throughput, cost
fig39	gateway throughput by worker count and the raw-loop ceiling
fig40	single port versus multi-port direct and via an L4 proxy
fig41	aggregate throughput and memory as the port count grows to 100
fig10–fig15	expert routing accuracy, cost, capacity

All figures in this paper are rendered on a white background at 300 dpi. Figures fig1–fig9 in `report/figures/` are early dark-theme renders kept for provenance only.